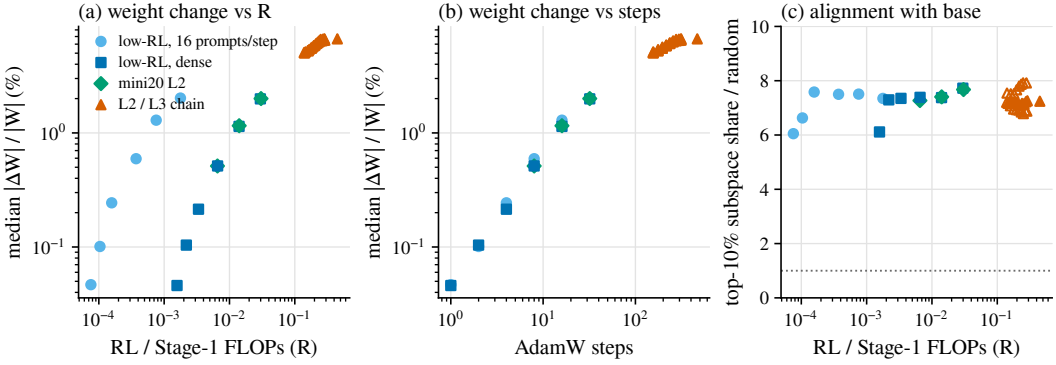
(a) weight change vs R
low-RL, 16 prompts/step
low-RL, dense
mini20 L2
L2 / L3 chain
median |ΔW| / |W| (%)
RL / Stage-1 FLOPs (R)
(b) weight change vs steps
median |ΔW| / |W| (%)
AdamW steps
(c) alignment with base
top-10% subspace share / random
RL / Stage-1 FLOPs (R)