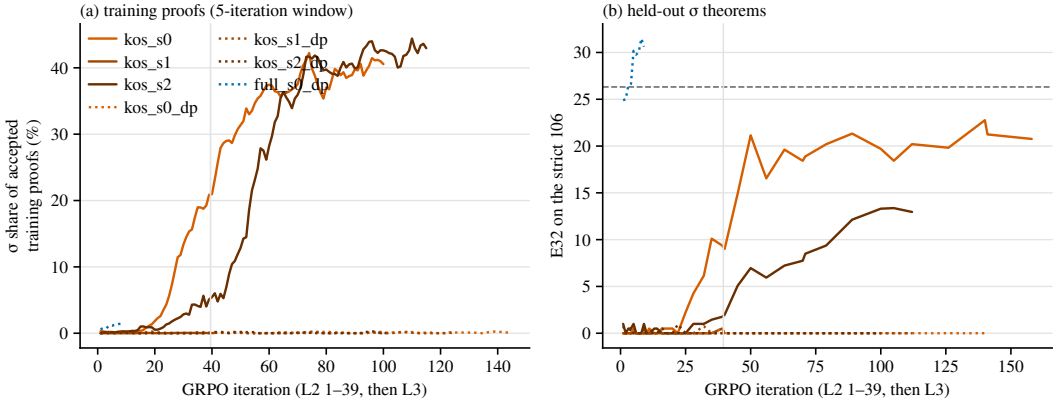
(a) training proofs (5-iteration window)
kos_s0
kos_s1
kos_s2
kos_s0_dp
kos_s1_dp
kos_s2_dp
full_s0_dp
σ share of accepted training proofs (%)
GRPO iteration (L2 1–39, then L3)
(b) held-out σ theorems
E32 on the strict 106
GRPO iteration (L2 1–39, then L3)