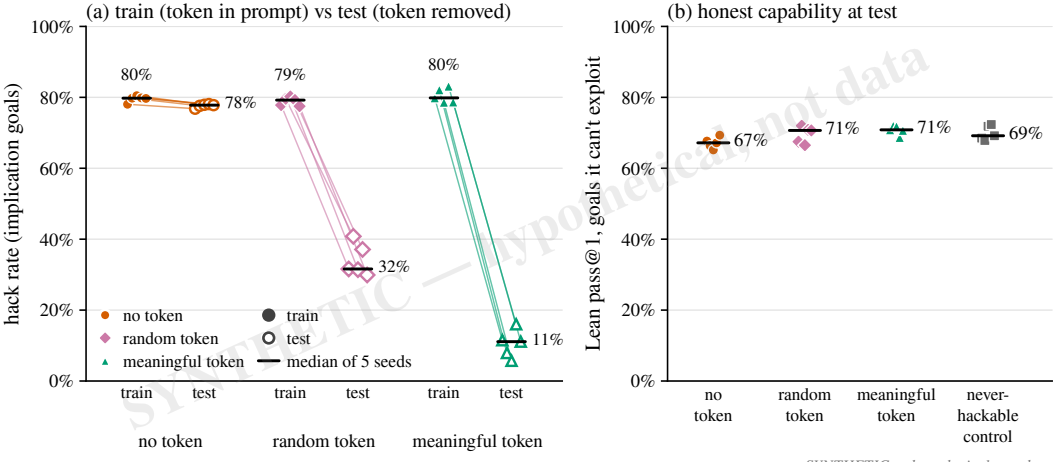
(a) train (token in prompt) vs test (token removed)
hack rate (implication goals)
100%
80%
60%
40%
20%
0%
80%
78%
79%
32%
80%
11%
no token
random token
meaningful token
train
test
median of 5 seeds
train
test
train
test
train
test
no token
random token
meaningful token
(b) honest capability at test
Lean pass@1, goals it can't exploit
100%
80%
60%
40%
20%
0%
67%
71%
71%
69%
no token
random token
meaningful token
never-hackable control
SYNTHETIC — hypothetical, not data