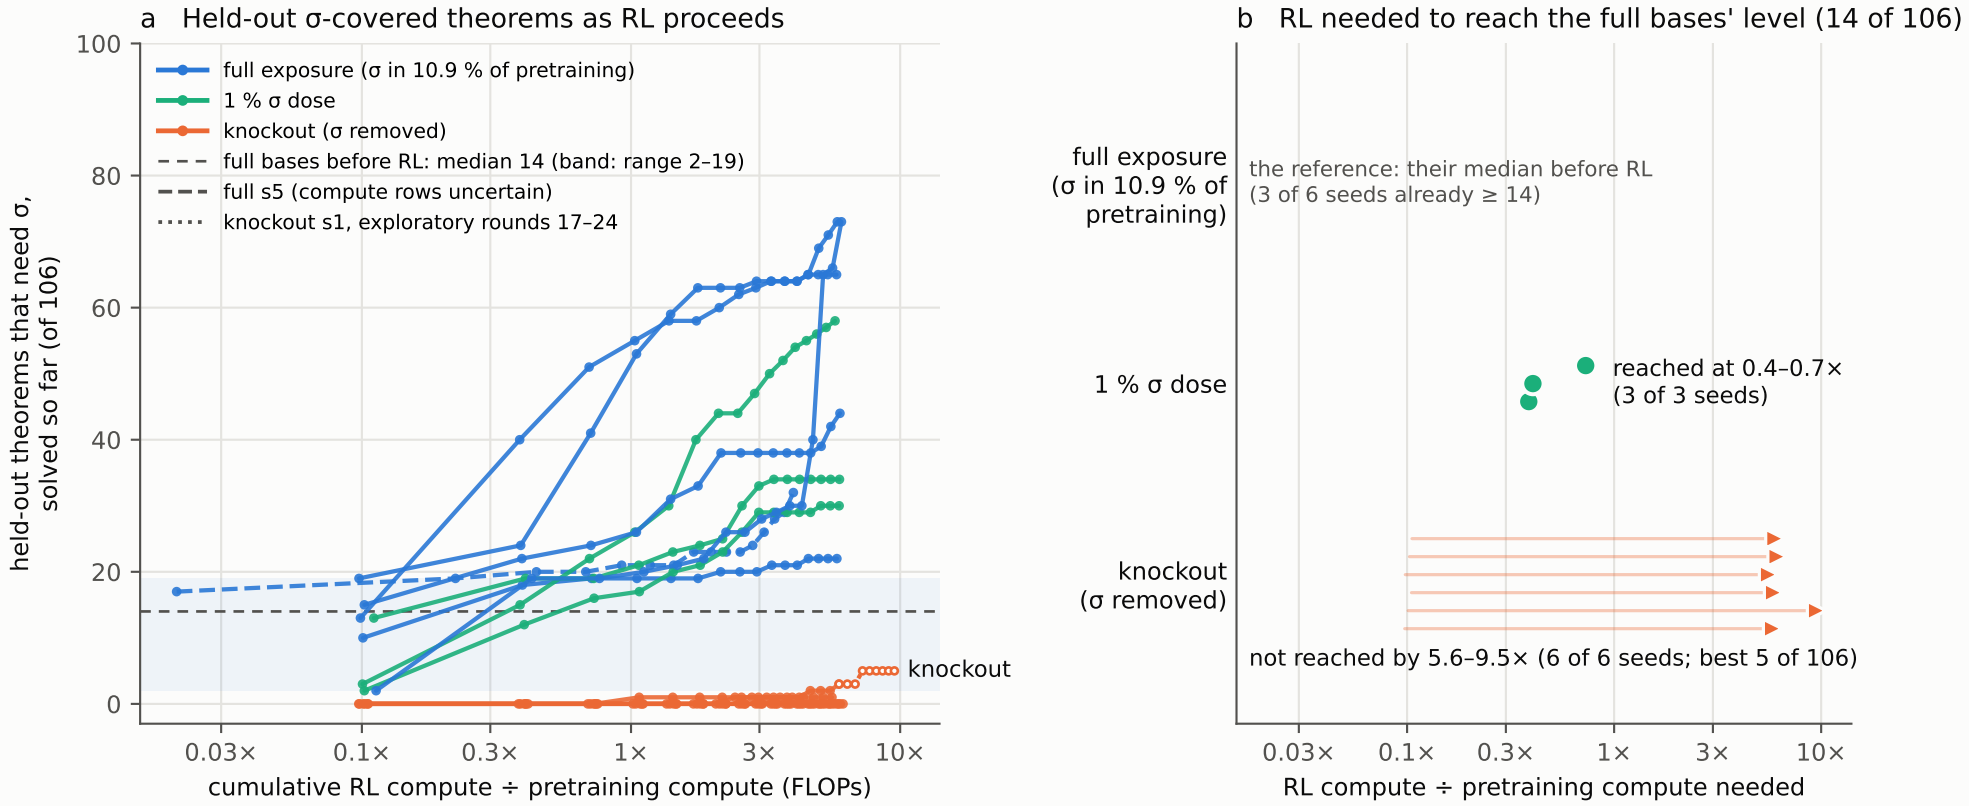

9.56M ALiBiGPT bases, T1 expert iteration, k 32 per round, judged by Lean (create-or-elicit). Held-out theorems are never trained on; a theorem counts once any round's sample solves it.
x: sampling (decoded and re-read state tokens) × 2N plus fine-tuning × 6N, including the 20,000 replayed pretraining records per round, over Stage-1's 6N·D (rl-compute-share's definitions; excluding replay lowers x by ≈ 14 %). Round 1 is the base's own samples, so it sits at ≈ 0.1×. Knockout values are lower bounds: those ladders stopped at round 16 (s1: 24).