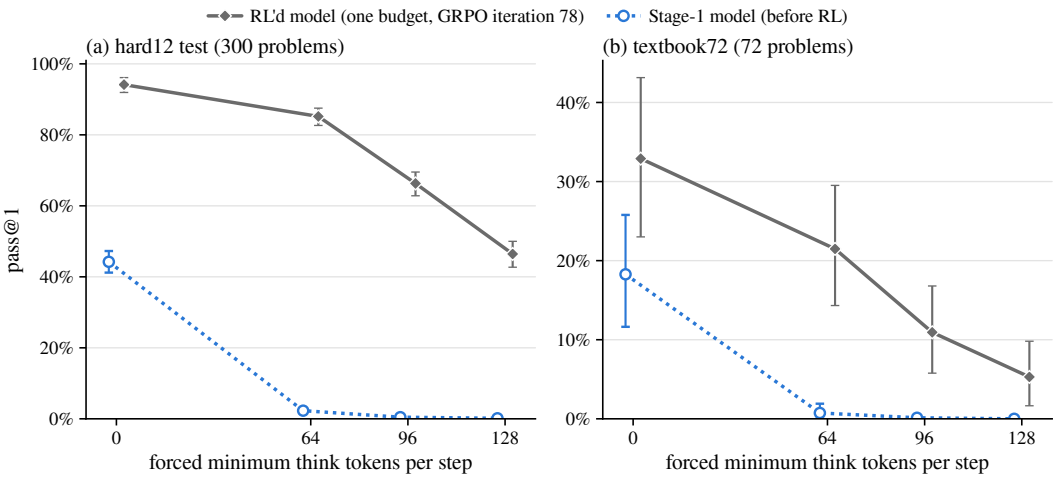
RL'd model (one budget, GRPO iteration 78)
Stage-1 model (before RL)
(a) hard12 test (300 problems)
(b) textbook72 (72 problems)
pass@1
100%
80%
60%
40%
20%
0%
40%
30%
20%
10%
0%
0
64
96
128
forced minimum think tokens per step
forced minimum think tokens per step