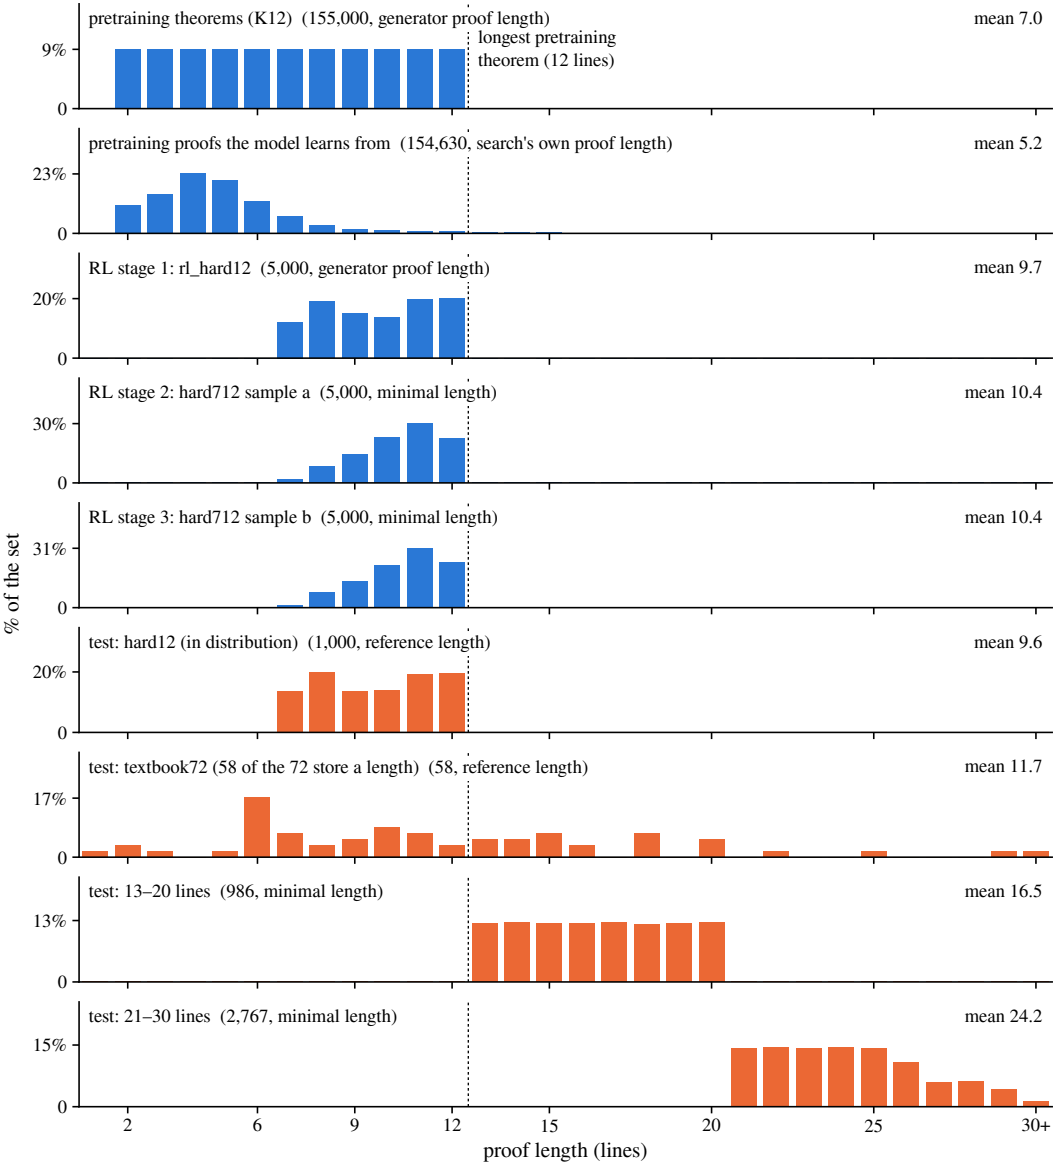
pretraining theorems (K12) (155,000, generator proof length)
mean 7.0
longest pretraining theorem (12 lines)
9%
0
pretraining proofs the model learns from (154,630, search's own proof length)
mean 5.2
23%
0
RL stage 1: rl_hard12 (5,000, generator proof length)
mean 9.7
20%
0
RL stage 2: hard712 sample a (5,000, minimal length)
mean 10.4
30%
0
RL stage 3: hard712 sample b (5,000, minimal length)
mean 10.4
31%
0
test: hard12 (in distribution) (1,000, reference length)
mean 9.6
20%
0
test: textbook72 (58 of the 72 store a length) (58, reference length)
mean 11.7
17%
0
test: 13–20 lines (986, minimal length)
mean 16.5
13%
0
test: 21–30 lines (2,767, minimal length)
mean 24.2
15%
0
% of the set
2
6
9
12
15
20
25
30+
proof length (lines)