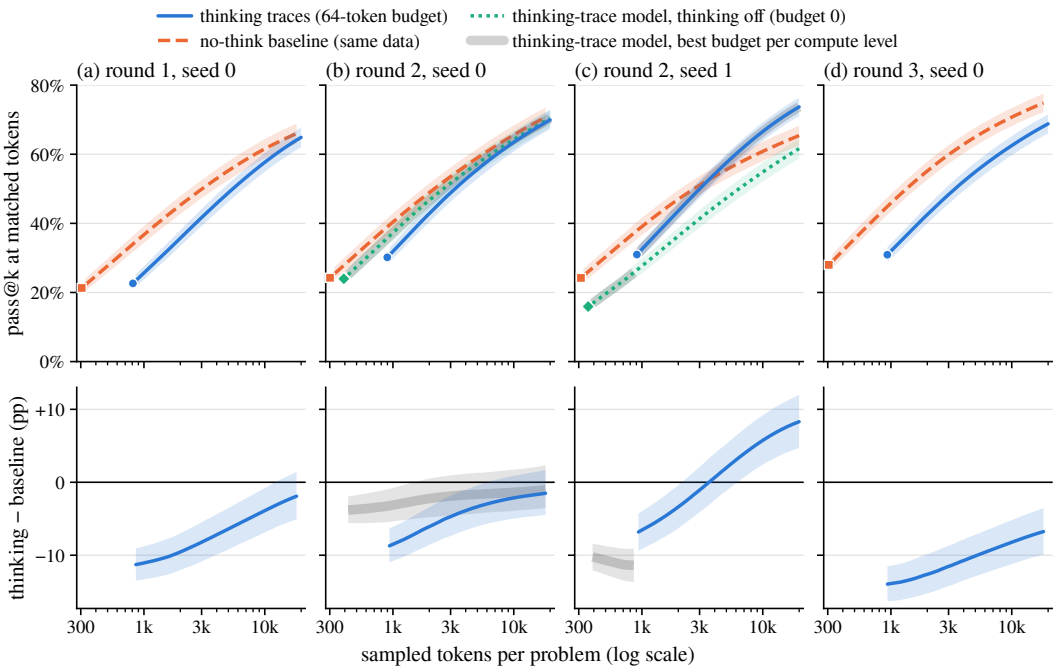
thinking traces (64-token budget)
no-think baseline (same data)
thinking-trace model, thinking off (budget 0)
thinking-trace model, best budget per compute level
(a) round 1, seed 0
(b) round 2, seed 0
(c) round 2, seed 1
(d) round 3, seed 0
pass@k at matched tokens
80%
60%
40%
20%
0%
thinking − baseline (pp)
+10
0
−10
300
1k
3k
10k
sampled tokens per problem (log scale)