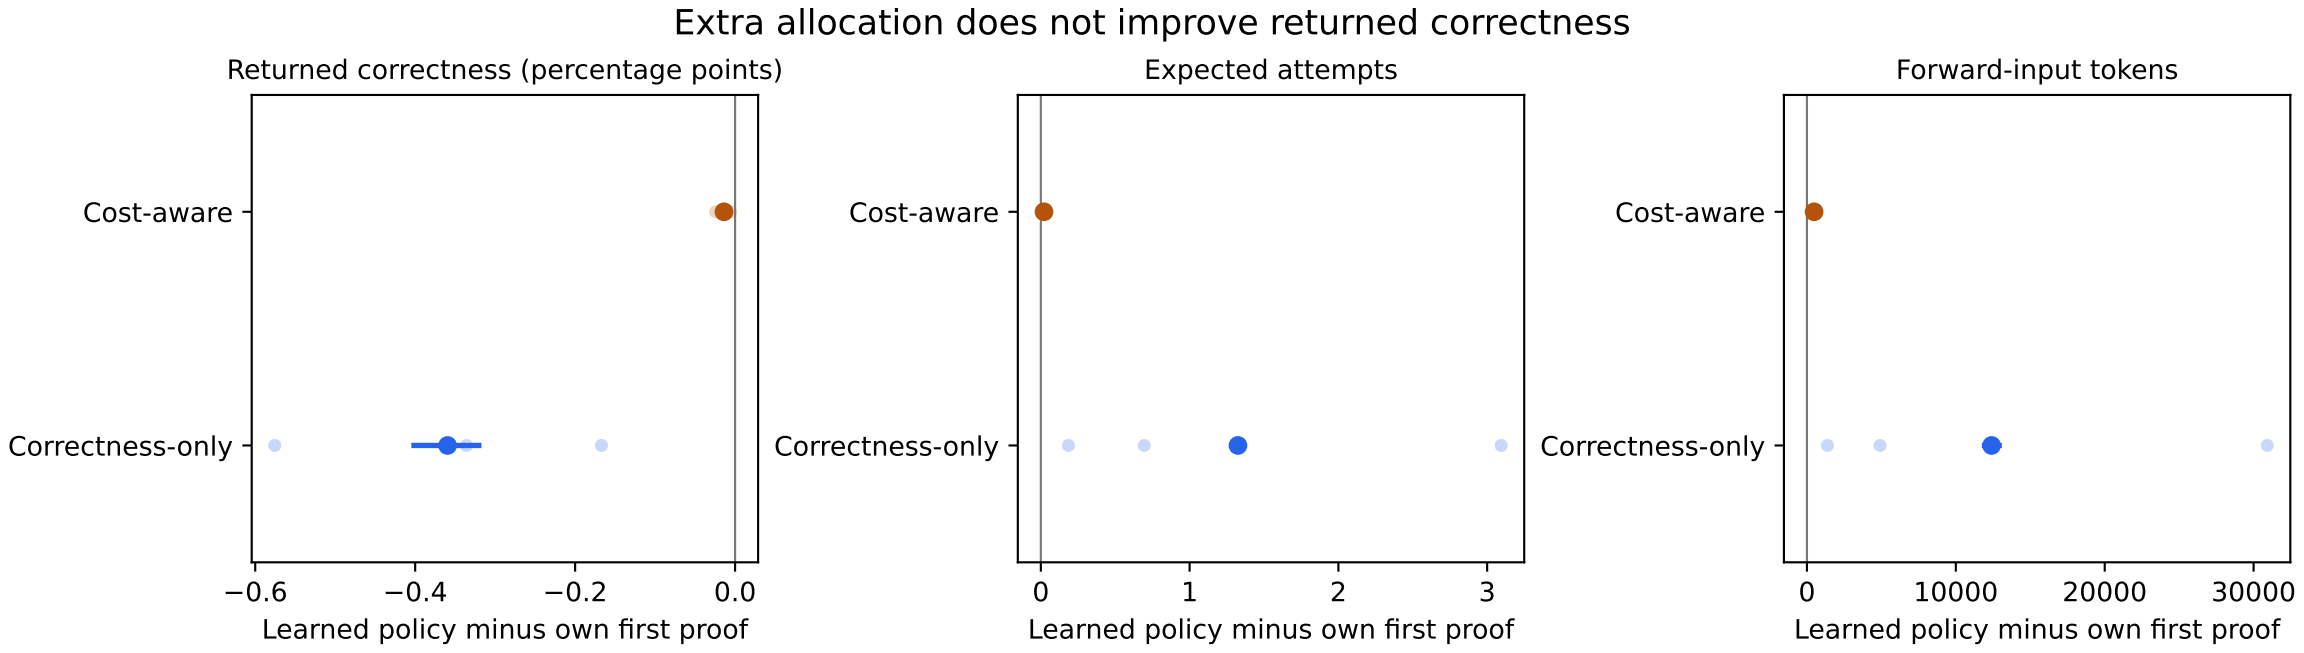
Extra allocation does not improve returned correctness
Returned correctness (percentage points)
Cost-aware
Correctness-only
−0.6
−0.4
−0.2
0.0
Learned policy minus own first proof
Expected attempts
Cost-aware
Correctness-only
0
1
2
3
Learned policy minus own first proof
Forward-input tokens
Cost-aware
Correctness-only
0
10000
20000
30000
Learned policy minus own first proof