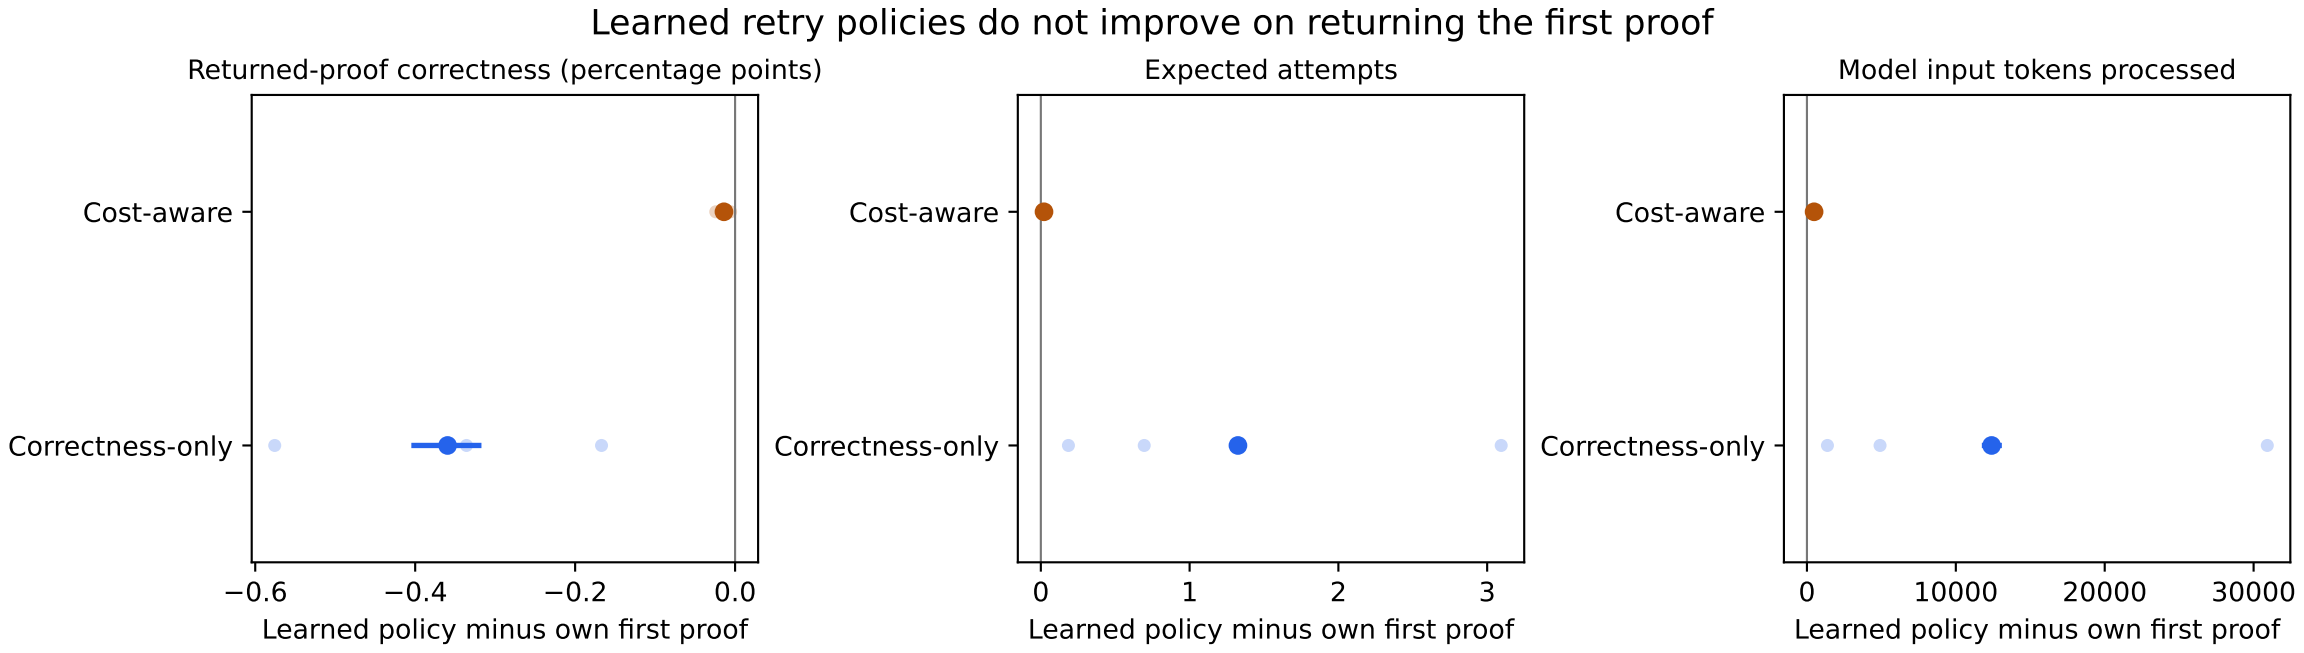
Learned retry policies do not improve on returning the first proof
Returned-proof correctness (percentage points)
Cost-aware
Correctness-only
−0.6
−0.4
−0.2
0.0
Learned policy minus own first proof
Expected attempts
Cost-aware
Correctness-only
0
1
2
3
Learned policy minus own first proof
Model input tokens processed
Cost-aware
Correctness-only
0
10000
20000
30000
Learned policy minus own first proof